


Binary Class Classification using Logistic Regression on Titanic Dataset

Titanic: Machine Learning from Disaster



Data Description

Overview

The data has been split into two groups:

- training set (train.csv)
- test set (test.csv)

The **training set** should be used to build your machine learning models. For the training set, we provide the outcome (also known as the “ground truth”) for each passenger. Your model will be based on “features” like passengers’ gender and class. You can also use feature engineering to create new features.

The **test set** should be used to see how well your model performs on unseen data. For the test set, we do not provide the ground truth for each passenger. It is your job to predict these outcomes. For each passenger in the test set, use the model you trained to predict whether or not they survived the sinking of the Titanic.

We also include gender_submission.csv, a set of predictions that assume all and only female passengers survive, as an example of what a submission file should look like.

Data Dictionary

Variable	Definition	Key
survival	Survival	0 = No, 1 = Yes
pclass	Ticket class	1 = 1st, 2 = 2nd, 3 = 3rd
sex	Sex	
Age	Age in years	
sibsp	# of siblings / spouses aboard the Titanic	
parch	# of parents / children aboard the Titanic	

Variable	Definition	Key
ticket	Ticket number	
fare	Passenger fare	
cabin	Cabin number	
embarked	Port of Embarkation	C = Cherbourg, Q = Queenstown, S = Southampton

Variable Notes

pclass: A proxy for socio-economic status (SES)

- 1st = Upper
- 2nd = Middle
- 3rd = Lower

age: Age is fractional if less than 1. If the age is estimated, is it in the form of xx.5

sibsp: The dataset defines family relations in this way...

- Sibling = brother, sister, stepbrother, stepsister
- Spouse = husband, wife (mistresses and fiancés were ignored)

parch: The dataset defines family relations in this way...

- Parent = mother, father
- Child = daughter, son, stepdaughter, stepson
- Some children travelled only with a nanny, therefore parch=0 for them.

In [1]:



```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import numpy as np
4 import seaborn as sns
```

In [2]:



```
1 df = pd.read_csv('datasets/Titanic.csv')
2 df.head()
```

Out[2]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	

In [3]:



```
1 df.shape
```

Out[3]:

(891, 12)

In [4]:



```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
PassengerId    891 non-null int64
Survived        891 non-null int64
Pclass         891 non-null int64
Name           891 non-null object
Sex            891 non-null object
Age           714 non-null float64
SibSp          891 non-null int64
Parch          891 non-null int64
Ticket         891 non-null object
Fare           891 non-null float64
Cabin          204 non-null object
Embarked       889 non-null object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.6+ KB
```

In [5]:



```
1 df.isnull().sum()
```

Out[5]:


```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            177
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          687
Embarked        2
dtype: int64
```

In [6]:

```
1 df['Sex'].value_counts().plot(kind = 'bar',color = 'g')
```

Out[6]:

<matplotlib.axes._subplots.AxesSubplot at 0x284daf22208>



In [7]:

```
1 df['Embarked'].value_counts()
```

Out[7]:

```
S    644
C    168
Q     77
Name: Embarked, dtype: int64
```

In [8]:

```
1 df['Embarked'].fillna('S',inplace = True)
```

In [9]:

```
1 df['Embarked'].value_counts()
```

Out[9]:


```
S    646
C    168
Q     77
Name: Embarked, dtype: int64
```

In [10]:

```
1 sns.barplot(df['Embarked'].value_counts().index,df['Embarked'].value_counts())
```

Out[10]:

<matplotlib.axes._subplots.AxesSubplot at 0x284db6633c8>



In [11]:

```
1 df.columns
```

Out[11]:

```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
      'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
      dtype='object')
```

In [12]:

```
1 df['Sex'] = pd.get_dummies(df['Sex'],drop_first=True)
```

In [13]:



```
1 df.head()
```

Out[13]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cal
0	1	0	3	Braund, Mr. Owen Harris	1	22.0	1	0	A/5 21171	7.2500	N
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	0	38.0	1	0	PC 17599	71.2833	C
2	3	1	3	Heikkinen, Miss. Laina	0	26.0	0	0	STON/O2. 3101282	7.9250	N
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	0	35.0	1	0	113803	53.1000	C1
4	5	0	3	Allen, Mr. William Henry	1	35.0	0	0	373450	8.0500	N



In [14]:



```
1 from sklearn.preprocessing import LabelEncoder
2
3 lb = LabelEncoder()
4
5 lb.fit_transform(df['Sex'])
```

Out[14]:

```
array([1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1,
       0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0,
       0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1,
       0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1,
       0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0,
       1, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1,
       0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1,
       1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1,
       1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1,
       0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0,
       1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1,
       1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0,
       1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0,
       0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1,
       1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1,
       0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1,
       0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0,
       1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1,
       0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0,
       1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1,
       1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1,
       0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0,
       1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0,
       1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1,
       0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0,
       0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1,
       0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1], dtype=int64)
```


In [15]:

```
1 df['Embarked'] = lb.fit_transform(df['Embarked'])
2 df.head()
```

Out[15]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	1	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	0	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	0	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	0	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	1	35.0	0	0	373450	8.0500

In [16]:

```
1 df.columns
```

Out[16]:

```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',  
      'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],  
      dtype='object')
```

In [17]:

```
1 req = df[['Survived', 'Pclass', 'Sex', 'Age', 'SibSp',  
2         'Parch', 'Fare', 'Embarked']]  
3  
4 cor = req.corr()
```

In [18]:

```
1 cor
```

Out[18]:


	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarke
Survived	1.000000	-0.338481	-0.543351	-0.077221	-0.035322	0.081629	0.257307	-0.16767
Pclass	-0.338481	1.000000	0.131900	-0.369226	0.083081	0.018443	-0.549500	0.16209
Sex	-0.543351	0.131900	1.000000	0.093254	-0.114631	-0.245489	-0.182333	0.10826
Age	-0.077221	-0.369226	0.093254	1.000000	-0.308247	-0.189119	0.096067	-0.03039
SibSp	-0.035322	0.083081	-0.114631	-0.308247	1.000000	0.414838	0.159651	0.06823
Parch	0.081629	0.018443	-0.245489	-0.189119	0.414838	1.000000	0.216225	0.03979
Fare	0.257307	-0.549500	-0.182333	0.096067	0.159651	0.216225	1.000000	-0.22471
Embarked	-0.167675	0.162098	0.108262	-0.030394	0.068230	0.039798	-0.224719	1.00000

In [19]:

```
1 plt.figure(figsize=(10,7))
2 sns.heatmap(cor,annot=True)
```

Out[19]:

<matplotlib.axes._subplots.AxesSubplot at 0x284db6cc780>



By observing the above correlation plot we can clearly state that there is no correlation between **Age** and other columns so we will fill the age with mean value.

In [20]:

```
1 req['Age'].fillna(req['Age'].mean(),inplace = True)
```

C:\Users\Jesus\Anaconda3\lib\site-packages\pandas\core\generic.py:6130: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

```
self._update_inplace(new_data)
```

In [21]:

```
1 req.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 8 columns):
Survived      891 non-null int64
Pclass        891 non-null int64
Sex           891 non-null uint8
Age           891 non-null float64
SibSp         891 non-null int64
Parch         891 non-null int64
Fare          891 non-null float64
Embarked      891 non-null int32
dtypes: float64(2), int32(1), int64(4), uint8(1)
memory usage: 46.2 KB
```

In [22]:

```
1 df.duplicated().sum()
```

Out[22]:

0

In [23]:



```
1 req[req.duplicated()]
```

Out[23]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
47	1	3	0	29.699118	0	0	7.7500	1
76	0	3	1	29.699118	0	0	7.8958	2
77	0	3	1	29.699118	0	0	8.0500	2
87	0	3	1	29.699118	0	0	8.0500	2
95	0	3	1	29.699118	0	0	8.0500	2
101	0	3	1	29.699118	0	0	7.8958	2
121	0	3	1	29.699118	0	0	8.0500	2
133	1	2	0	29.000000	1	0	26.0000	2
173	0	3	1	21.000000	0	0	7.9250	2
196	0	3	1	29.699118	0	0	7.7500	1
198	1	3	0	29.699118	0	0	7.7500	1
201	0	3	1	29.699118	8	2	69.5500	2
213	0	2	1	30.000000	0	0	13.0000	2
223	0	3	1	29.699118	0	0	7.8958	2
241	1	3	0	29.699118	1	0	15.5000	1
260	0	3	1	29.699118	0	0	7.7500	1
274	1	3	0	29.699118	0	0	7.7500	1
295	0	1	1	29.699118	0	0	27.7208	0
300	1	3	0	29.699118	0	0	7.7500	1
304	0	3	1	29.699118	0	0	8.0500	2
313	0	3	1	28.000000	0	0	7.8958	2
320	0	3	1	22.000000	0	0	7.2500	2
324	0	3	1	29.699118	8	2	69.5500	2
335	0	3	1	29.699118	0	0	7.8958	2
343	0	2	1	25.000000	0	0	13.0000	2
354	0	3	1	29.699118	0	0	7.2250	0
355	0	3	1	28.000000	0	0	9.5000	2
358	1	3	0	29.699118	0	0	7.8792	1
359	1	3	0	29.699118	0	0	7.8792	1
364	0	3	1	29.699118	1	0	15.5000	1
...
692	1	3	1	29.699118	0	0	56.4958	2
696	0	3	1	44.000000	0	0	8.0500	2
709	1	3	1	29.699118	1	1	15.2458	0

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
732	0	2	1	29.699118	0	0	0.0000	2
733	0	2	1	23.000000	0	0	13.0000	2
734	0	2	1	23.000000	0	0	13.0000	2
738	0	3	1	29.699118	0	0	7.8958	2
739	0	3	1	29.699118	0	0	7.8958	2
757	0	2	1	18.000000	0	0	11.5000	2
758	0	3	1	34.000000	0	0	8.0500	2
760	0	3	1	29.699118	0	0	14.5000	2
773	0	3	1	29.699118	0	0	7.2250	0
776	0	3	1	29.699118	0	0	7.7500	1
790	0	3	1	29.699118	0	0	7.7500	1
792	0	3	0	29.699118	8	2	69.5500	2
800	0	2	1	34.000000	0	0	13.0000	2
808	0	2	1	39.000000	0	0	13.0000	2
815	0	1	1	29.699118	0	0	0.0000	2
832	0	3	1	29.699118	0	0	7.2292	0
837	0	3	1	29.699118	0	0	8.0500	2
838	1	3	1	32.000000	0	0	56.4958	2
844	0	3	1	17.000000	0	0	8.6625	2
846	0	3	1	29.699118	8	2	69.5500	2
859	0	3	1	29.699118	0	0	7.2292	0
863	0	3	0	29.699118	8	2	69.5500	2
870	0	3	1	26.000000	0	0	7.8958	2
877	0	3	1	19.000000	0	0	7.8958	2
878	0	3	1	29.699118	0	0	7.8958	2
884	0	3	1	25.000000	0	0	7.0500	2
886	0	2	1	27.000000	0	0	13.0000	2

111 rows × 8 columns

In [24]:



```
1 Y = req['Survived']
2 X = req.drop('Survived',axis = 1)
```

In [25]:



```
1 from sklearn.model_selection import train_test_split
2
3 X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.3, random_state = 4)
```

In [26]:



```
1 X.head()
```

Out[26]:

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	1	22.0	1	0	7.2500	2
1	1	0	38.0	1	0	71.2833	0
2	3	0	26.0	0	0	7.9250	2
3	1	0	35.0	1	0	53.1000	2
4	3	1	35.0	0	0	8.0500	2

In [27]:



```
1 X_train.head()
```

Out[27]:

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
445	1	1	4.000000	0	2	81.8583	2
650	3	1	29.699118	0	0	7.8958	2
172	3	0	1.000000	1	1	11.1333	2
450	2	1	36.000000	1	2	27.7500	2
314	2	1	43.000000	1	1	26.2500	2

In [28]:



```
1 X_train.head()
```

Out[28]:

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
445	1	1	4.000000	0	2	81.8583	2
650	3	1	29.699118	0	0	7.8958	2
172	3	0	1.000000	1	1	11.1333	2
450	2	1	36.000000	1	2	27.7500	2
314	2	1	43.000000	1	1	26.2500	2

In [29]:



```
1 from sklearn.linear_model import LogisticRegression
2
3 logreg = LogisticRegression()
4
5 logreg.fit(X,Y)
```

C:\Users\Jesus\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
FutureWarning)

Out[29]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False)
```

In [30]:



```
1 logreg.predict([[1,1,4.000000,0,2,81.8583,2]])
```

Out[30]:

```
array([1], dtype=int64)
```

In [31]:



```
1 Y_train.head(1)
```

Out[31]:

```
445    1
Name: Survived, dtype: int64
```

In [32]:



```
1 logreg.predict_proba([[1,1,4.000000,0,2,81.8583,2]])
```

Out[32]:

```
array([[0.3874428, 0.6125572]])
```

In [33]:



```
1 Y_pred = logreg.predict(X_test)
```

In [34]:



```
1 Y_pred
```

Out[34]:

```
array([0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
       1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
       1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1,
       0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
       1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0,
       0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
       0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0,
       0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0,
       1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0,
       0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1,
       0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0,
       0, 0, 0], dtype=int64)
```

In [35]:



```
1 from sklearn.metrics import confusion_matrix
2 con = confusion_matrix(Y_pred,Y_test)
3 con
```

Out[35]:

```
array([[138,  31],
       [ 19,  80]], dtype=int64)
```

In [36]:



```
1 Y_test.shape
```

Out[36]:

```
(268,)
```

In [37]:



```
1 con.sum()
```

Out[37]:

```
268
```

In [38]:



```
1 from sklearn.metrics import accuracy_score
2
3 accuracy_score(Y_pred,Y_test)
```

Out[38]:

```
0.8134328358208955
```


In [39]:



```
1 from sklearn.metrics import classification_report
2
3 cp = classification_report(Y_pred,Y_test)
4
5 print(cp)
```

	precision	recall	f1-score	support
0	0.88	0.82	0.85	169
1	0.72	0.81	0.76	99
micro avg	0.81	0.81	0.81	268
macro avg	0.80	0.81	0.80	268
weighted avg	0.82	0.81	0.82	268

Binary Class Classification using Logistic Regression on Breast Cancer Dataset

In [40]:



```
1 from sklearn.datasets import load_breast_cancer
2
3 data = load_breast_cancer()
4
5 df.keys()
```

Out[40]:

```
Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
      'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
      dtype='object')
```

In [41]:



```
1 data.target_names
```

Out[41]:

```
array(['malignant', 'benign'], dtype='<U9')
```

In [42]:



```
1 data.feature_names
```

Out[42]:

```
array(['mean radius', 'mean texture', 'mean perimeter', 'mean area',  
      'mean smoothness', 'mean compactness', 'mean concavity',  
      'mean concave points', 'mean symmetry', 'mean fractal dimension',  
      'radius error', 'texture error', 'perimeter error', 'area error',  
      'smoothness error', 'compactness error', 'concavity error',  
      'concave points error', 'symmetry error',  
      'fractal dimension error', 'worst radius', 'worst texture',  
      'worst perimeter', 'worst area', 'worst smoothness',  
      'worst compactness', 'worst concavity', 'worst concave points',  
      'worst symmetry', 'worst fractal dimension'], dtype='<U23')
```

In [44]:



```
1 print(data.DESCR)
```

.. _breast_cancer_dataset:

Breast cancer wisconsin (diagnostic) dataset

****Data Set Characteristics:****

- :Number of Instances: 569
- :Number of Attributes: 30 numeric, predictive attributes and the class
- :Attribute Information:
 - radius (mean of distances from center to points on the perimeter)
 - texture (standard deviation of gray-scale values)
 - perimeter
 - area
 - smoothness (local variation in radius lengths)
 - compactness (perimeter² / area - 1.0)
 - concavity (severity of concave portions of the contour)
 - concave points (number of concave portions of the contour)
 - symmetry
 - fractal dimension ("coastline approximation" - 1)

The mean, standard error, and "worst" or largest (mean of the three largest values) of these features were computed for each image, resulting in 30 features. For instance, field 3 is Mean Radius, fie

ld

13 is Radius SE, field 23 is Worst Radius.

- class:
 - WDBC-Malignant
 - WDBC-Benign

:Summary Statistics:

=====	=====	=====
	Min	Max
=====	=====	=====
radius (mean):	6.981	28.11
texture (mean):	9.71	39.28
perimeter (mean):	43.79	188.5
area (mean):	143.5	2501.0
smoothness (mean):	0.053	0.163
compactness (mean):	0.019	0.345
concavity (mean):	0.0	0.427
concave points (mean):	0.0	0.201
symmetry (mean):	0.106	0.304
fractal dimension (mean):	0.05	0.097
radius (standard error):	0.112	2.873
texture (standard error):	0.36	4.885
perimeter (standard error):	0.757	21.98
area (standard error):	6.802	542.2
smoothness (standard error):	0.002	0.031
compactness (standard error):	0.002	0.135
concavity (standard error):	0.0	0.396
concave points (standard error):	0.0	0.053

symmetry (standard error):	0.008	0.079
fractal dimension (standard error):	0.001	0.03
radius (worst):	7.93	36.04
texture (worst):	12.02	49.54
perimeter (worst):	50.41	251.2
area (worst):	185.2	4254.0
smoothness (worst):	0.071	0.223
compactness (worst):	0.027	1.058
concavity (worst):	0.0	1.252
concave points (worst):	0.0	0.291
symmetry (worst):	0.156	0.664
fractal dimension (worst):	0.055	0.208
=====	=====	=====

:Missing Attribute Values: None

:Class Distribution: 212 - Malignant, 357 - Benign

:Creator: Dr. William H. Wolberg, W. Nick Street, Olvi L. Mangasarian

:Donor: Nick Street

:Date: November, 1995

This is a copy of UCI ML Breast Cancer Wisconsin (Diagnostic) datasets.

<https://goo.gl/U2Uwz2> (<https://goo.gl/U2Uwz2>)

Features are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image.

Separating plane described above was obtained using Multisurface Method-Tree (MSM-T) [K. P. Bennett, "Decision Tree Construction Via Linear Programming." Proceedings of the 4th Midwest Artificial Intelligence and Cognitive Science Society, pp. 97-101, 1992], a classification method which uses linear programming to construct a decision tree. Relevant features were selected using an exhaustive search in the space of 1-4 features and 1-3 separating planes.

The actual linear program used to obtain the separating plane in the 3-dimensional space is that described in:

[K. P. Bennett and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets", Optimization Methods and Software 1, 1992, 23-34].

This database is also available through the UW CS ftp server:

```
ftp ftp.cs.wisc.edu
cd math-prog/cpo-dataset/machine-learn/WDBC/
```

.. topic:: References

- W.N. Street, W.H. Wolberg and O.L. Mangasarian. Nuclear feature extraction for breast tumor diagnosis. IS&T/SPIE 1993 International Symposium on Electronic Imaging: Science and Technology, volume 1905, pages 861-870, San Jose, CA, 1993.
- O.L. Mangasarian, W.N. Street and W.H. Wolberg. Breast cancer diagnosis and prognosis via linear programming. Operations Research, 43(4), pages 570

-577,
July-August 1995.
- W.H. Wolberg, W.N. Street, and O.L. Mangasarian. Machine learning techniques
to diagnose breast cancer from fine-needle aspirates. Cancer Letters 77
(1994)
163-171.

In [45]:

```
1 data.data
```

Out[45]:

```
array([[1.799e+01, 1.038e+01, 1.228e+02, ..., 2.654e-01, 4.601e-01,
        1.189e-01],
       [2.057e+01, 1.777e+01, 1.329e+02, ..., 1.860e-01, 2.750e-01,
        8.902e-02],
       [1.969e+01, 2.125e+01, 1.300e+02, ..., 2.430e-01, 3.613e-01,
        8.758e-02],
       ...,
       [1.660e+01, 2.808e+01, 1.083e+02, ..., 1.418e-01, 2.218e-01,
        7.820e-02],
       [2.060e+01, 2.933e+01, 1.401e+02, ..., 2.650e-01, 4.087e-01,
        1.240e-01],
       [7.760e+00, 2.454e+01, 4.792e+01, ..., 0.000e+00, 2.871e-01,
        7.039e-02]])
```

In [46]:

```
1 import pandas as pd
2
3 canc = pd.DataFrame(data.data,columns=data.feature_names)
4
5 canc.head()
```

Out[46]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809

5 rows × 30 columns

In [47]:

```
1 canc['output'] = data.target
```

In [48]:

```
1 canc.head()
```

Out[48]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809

5 rows × 31 columns

In [49]:

```
1 X = canc.drop('output',axis = 1)
2 Y = canc['output']
```

In [50]:

```
1 from sklearn.model_selection import train_test_split
2
3 X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.3, random_state = 4)
```

In [51]:

```
1 from sklearn.linear_model import LogisticRegression
2
3 lgc = LogisticRegression()
4
5 lgc.fit(X_train,Y_train)
6
7 Y_pred = lgc.predict(X_test)
```

C:\Users\Jesus\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specifiy a solver to silence this warning.
FutureWarning)

In [52]:



```
1 from sklearn.metrics import accuracy_score, confusion_matrix
2
3 print('Confusion Matrix:\n', confusion_matrix(Y_test, Y_pred))
4
5 print('accuracy_score:', accuracy_score(Y_test, Y_pred))
```

Confusion Matrix:

```
[[ 59   4]
 [  2 106]]
```

accuracy_score: 0.9649122807017544