

Multiple Linear Regression Model

A linear regression model that contains more than one predictor variable is called a multiple linear regression model. The following model is a multiple linear regression model with two predictor variables, X_1 and X_2 . then

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon$$

For more description [Click Here](http://reliawiki.org/index.php/Multiple_Linear_Regression_Analysis)
(http://reliawiki.org/index.php/Multiple_Linear_Regression_Analysis)

[Datasets Link \(tiny.cc/apssdc-datasets\)](http://tiny.cc/apssdc-datasets)

In [1]:

```
import pandas as pd
```

In [2]:

```
fuel = pd.read_csv('FuelConsumptionCo2.csv')  
fuel.head()
```

Out[2]:

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINE SIZE	CYLINDERS	TRANSMISSION
0	2014	ACURA	ILX	COMPACT	2.0	4	AS
1	2014	ACURA	ILX	COMPACT	2.4	4	M
2	2014	ACURA	ILX HYBRID	COMPACT	1.5	4	AV
3	2014	ACURA	MDX 4WD	SUV - SMALL	3.5	6	AS
4	2014	ACURA	RDX AWD	SUV - SMALL	3.5	6	AS

In [3]:

```
fuel.shape
```

Out[3]:

(1067, 13)

In [4]:

```
fuel.columns
```

Out[4]:

```
Index(['MODELYEAR', 'MAKE', 'MODEL', 'VEHICLECLASS', 'ENGINESIZE', 'CYLINDERS',  
      'TRANSMISSION', 'FUELTYPE', 'FUELCONSUMPTION_CITY',  
      'FUELCONSUMPTION_HWY', 'FUELCONSUMPTION_COMB',  
      'FUELCONSUMPTION_COMB_MPG', 'CO2EMISSIONS'],  
      dtype='object')
```

Understanding the Data

FuelConsumption.csv :

- **MODELYEAR** e.g. 2014
- **MAKE** e.g. Acura
- **MODEL** e.g. ILX
- **VEHICLE CLASS** e.g. SUV
- **ENGINE SIZE** e.g. 4.7
- **CYLINDERS** e.g 6
- **TRANSMISSION** e.g. A6
- **FUEL CONSUMPTION in CITY(L/100 km)** e.g. 9.9
- **FUEL CONSUMPTION in HWY (L/100 km)** e.g. 8.9
- **FUEL CONSUMPTION COMB (L/100 km)** e.g. 9.2
- **CO2 EMISSIONS (g/km)** e.g. 182 --> low --> 0

In [5]:

```
len(fuel['MAKE'].value_counts())
```

Out[5]:

39

In [6]:

```
fuel['MAKE'].value_counts()
```

Out[6]:

FORD	90
CHEVROLET	86
BMW	64
MERCEDES-BENZ	59
GMC	49
TOYOTA	49
AUDI	49
PORSCHE	44
VOLKSWAGEN	42
DODGE	39
MINI	36
NISSAN	33
KIA	33
CADILLAC	32
JEEP	31
MAZDA	27
HYUNDAI	24
SUBARU	23
LEXUS	22
JAGUAR	22
HONDA	21
INFINITI	21
LAND ROVER	19
CHRYSLER	19
MITSUBISHI	16
BUICK	16
RAM	13
ACURA	12
VOLVO	11
LINCOLN	11
FIAT	10
SCION	9
BENTLEY	8
ROLLS-ROYCE	7
ASTON MARTIN	7
MASERATI	6
LAMBORGHINI	3
SMART	2
SRT	2

Name: MAKE, dtype: int64

In [7]:

```
features = fuel[['ENGINE_SIZE', 'CYLINDERS', 'FUEL_CONSUMPTION_COMB', 'CO2_EMISSIONS']]
features.head()
```

Out[7]:

	ENGINE_SIZE	CYLINDERS	FUEL_CONSUMPTION_COMB	CO2_EMISSIONS
0	2.0	4	8.5	196
1	2.4	4	9.6	221
2	1.5	4	5.9	136
3	3.5	6	11.1	255
4	3.5	6	10.6	244

In [8]:

```
import seaborn as sns
```

```
sns.pairplot(data = features)
```

In [9]:

```
X = features[['ENGINE_SIZE', 'FUEL_CONSUMPTION_COMB']]
Y = features[['CO2_EMISSIONS']]
```

In [10]:

```
from sklearn.linear_model import LinearRegression
```

In [11]:

```
reg = LinearRegression() # Regression Object
pred = reg.fit(X,Y)
Y_pred = reg.predict(X)
```

In [12]:

```
reg.score(X,Y)
```

Out[12]:

```
0.8582144778981454
```

In [13]:

```
reg.predict([[2.0,8.5]])
```

Out[13]:

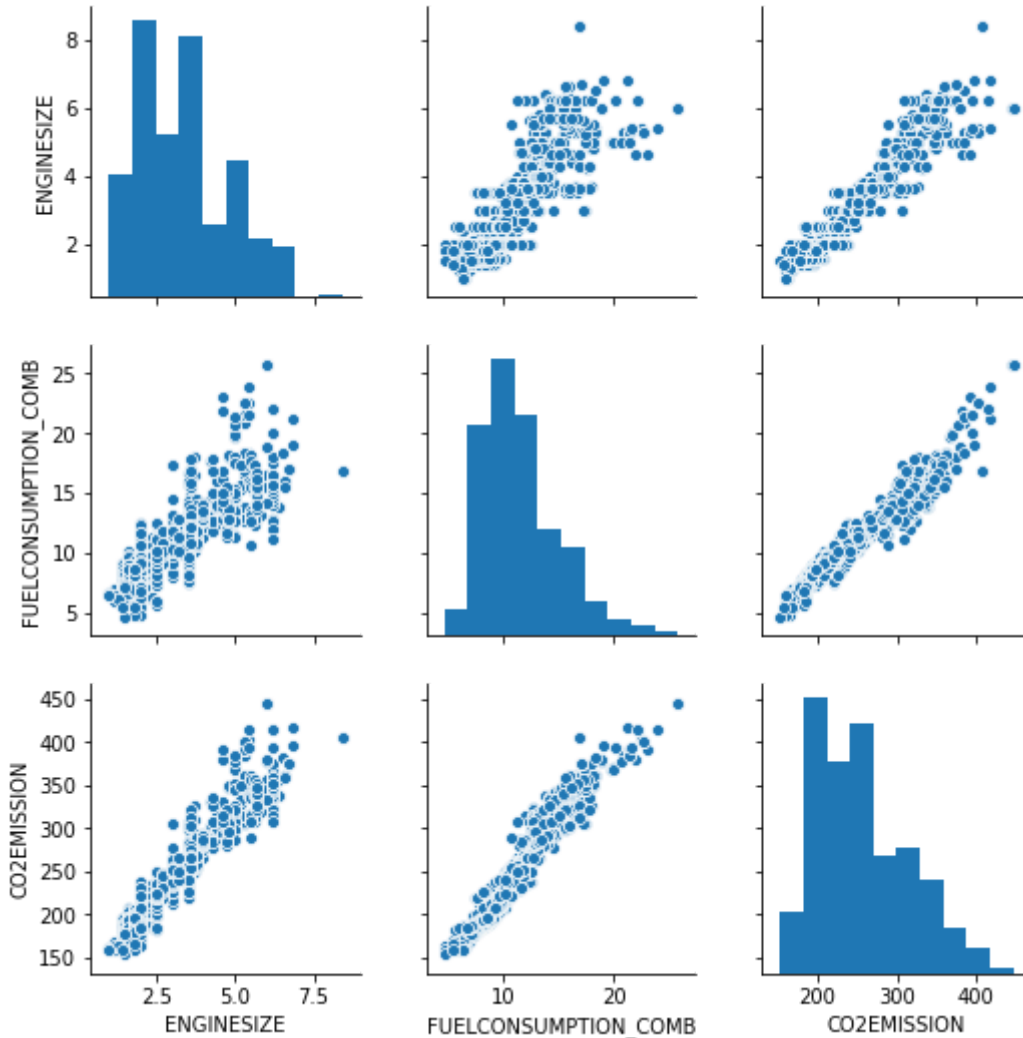
```
array([[200.00402631]])
```

In [14]:

```
df = pd.DataFrame(X)
df['CO2EMISSION'] = pd.DataFrame(Y_pred)
sns.pairplot(df)
```

Out[14]:

<seaborn.axisgrid.PairGrid at 0x18d42f396a0>



In [15]:

```
reg.coef_
```

Out[15]:

```
array([[19.49628353,  9.72995852]])
```

In [16]:

```
reg.intercept_
```

Out[16]:

```
array([78.30681186])
```

In [17]:

```
from sklearn.model_selection import train_test_split
```

In [18]:

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.3, random_state = 7)
```

In [19]:

```
reg1 = LinearRegression()  
reg1.fit(X_train,Y_train)  
Y_pred = reg1.predict(X_test)  
reg1.score(X_test,Y_test)
```

Out[19]:

0.8721283450220715

In [20]:

```
Y_pred = reg1.predict(X_test)
```

In [21]:

```
reg1.score(X_test,Y_test)
```

Out[21]:

0.8721283450220715